

# Architectural Analysis and Validation of a Gymnasium-Based Order-Driven Market Simulator

[Mudil Goel]

December 29, 2025

## 1 Market Structure & Matching Logic

The core of the simulator is modeled as an **Order-Driven Market**, a structural choice that aligns with the dominant design of modern global equity exchanges (e.g., NYSE, Nasdaq, Tokyo Stock Exchange). Unlike quote-driven markets, where designated market makers or dealers set buy and sell prices, an order-driven market relies on the interaction of limit orders submitted by market participants to determine the price.

### 1.1 The Limit Order Book (LOB)

The central data structure of the simulator is the `OrderBook` class. This class maintains two distinct, sorted lists: `bids` (buy orders) and `asks` (sell orders). In adherence to standard microstructure theory, the simulator treats liquidity as a scarce resource.

- **Bid Side:** Sorted in descending order. The highest price a participant is willing to pay is at the “top of the book.”
- **Ask Side:** Sorted in ascending order. The lowest price a participant is willing to accept is prioritized.

### 1.2 Matching Mechanism: Price-Time Priority

To resolve liquidity conflicts, the simulator implements a **Price-Time Priority (FIFO)** algorithm. This is the industry standard for equity markets, designed to incentivize liquidity provision by rewarding traders who improve the best available price or, in the event of a tie, those who arrive earliest.

In the Python implementation, this logic is encapsulated in the `match()` method. The engine continuously monitors the spread. When the inequality  $\text{BestBid} \geq \text{BestAsk}$  is satisfied, a trade occurs. The system executes the trade at the bid price (a simplified crossing rule), records the transaction on the “tape” (the `trades` list), and handles partial fills by decrementing the quantity of the matched orders. This mimics the continuous double auction mechanism found in real-world electronic communication networks (ECNs).

## 2 The Agent-Environment Loop

A critical architectural requirement for this project was compatibility with Reinforcement Learning (RL) workflows. To achieve this, the simulation is wrapped in a custom `MarketEnvironment` class that adheres to the **Gymnasium (formerly OpenAI Gym)** API standards.

### 2.1 Separation of Concerns

The architecture enforces a strict separation between the market mechanism and the trader.

- **The Environment (TradingEnv):** Acts as the exchange operator. It enforces market rules (e.g., rejecting trades when the book is empty), manages the passage of time, and calculates the results of actions.
- **The Agent:** Represents the trader. The agent submits actions (orders) without direct access to the internal logic of the matching engine.

### 2.2 Discrete-Time State Transitions

Markets in the real world are continuous, but for the purpose of algorithmic training, the simulator discretizes time into “steps.” Each call to the `step()` function represents a micro-moment where the agent assesses the market and submits an order. The state transition follows the Markov Decision Process (MDP) tuple  $(S, A, R, S')$ , where the agent observes state  $S$ , takes action  $A$  (Buy/Sell, Price, Qty), receives reward  $R$  (profit/loss), and the market transitions to a new state  $S'$ . This standardization allows for the seamless future integration of advanced RL algorithms such as PPO or DQN.

## 3 Information & State Observation

A defining characteristic of financial markets is **Information Asymmetry**. As noted by O’Hara in *Market Microstructure Theory*, traders almost never possess perfect information about the “true” value of an asset or the hidden intent of other participants.

### 3.1 Level 2 (L2) Market Data

To replicate this constraint, the simulator provides agents with a restricted view of the world known as “Level 2 Market Data.” The `get_l2_state()` function generates an observation vector containing only the prices and quantities of the top  $N$  bids and asks.

### 3.2 Implementation of Partial Observability

Crucially, the agent **cannot** see:

- The identity of other traders (preserving anonymity).
- Orders buried deep in the book (simulating “dark” or off-screen liquidity).
- The private valuation or “alpha” of other agents.

This design forces the agent to infer market dynamics solely from the public flow of orders, mimicking the real-world experience of a quantitative trader staring at a Bloomberg terminal or an order depth ladder.

## 4 Validation: The Zero-Intelligence Test

To verify the integrity of the market mechanism before introducing complex AI strategies, the system was subjected to a **Zero-Intelligence (ZI) Constraint Test**. This methodology is derived from the seminal paper by Gode and Sunder (1993), *Allocative Efficiency of Markets with Zero-Intelligence Traders*.

### 4.1 Methodology

The simulation was initialized with a `RandomAgent`—a trader programmed to submit orders with random sides (buy/sell), random price offsets ( $P_{last} \pm \delta$ ), and random quantities. This agent possesses no strategy, no memory, and no intent to maximize profit.

### 4.2 Results and Analysis

The simulation was run for 100 timesteps. The resulting price path exhibited the following characteristics:

- **Price Discovery:** Despite the random inputs, the transaction prices did not explode to infinity or collapse to zero. They remained within a bounded range (93.0 to 101.0), oscillating around the initial equilibrium price.
- **Non-Trivial Liquidity:** The presence of a fluctuating price path confirms that the `match()` logic successfully identified and executed crossing orders.

### 4.3 Conclusion

The successful execution of the Zero-Intelligence test confirms the “Allocative Efficiency” of the simulator. It proves that the market structure itself—the funneling of bids and asks through a priority engine—is sufficient to generate a coherent price path. The simulator is therefore validated as a robust environment for the next phase of research: training intelligent agents to exploit these structural dynamics.